

# **In-Class Assessment 2:**

## **Kubernetes Orchestration, Deployment and Cluster Management**

Name: Yin Yuang

Student ID: 202383930027

Class: Software Engineering Class 1

Introduction to Cloud Computing  
(Autumn,2025)

Nanjing University of Information Science and Technology, China  
Waterford Institute

## **1. Orchestration tools, such as Kubernetes, play a key role in server infrastructure for modern applications**

### **(a) Explain how these tools help manage and scale application servers**

Kubernetes helps manage and scale application servers through the following mechanisms:

1. **Automated Operations:** Automatically handles Pod deployment, scaling, failure recovery, and rolling updates.
2. **Resource Management:** Precisely controls CPU and memory allocation using Requests and Limits.
3. **Elastic Scaling:** Supports Horizontal Pod Autoscaler (HPA) to dynamically adjust replica count based on workload.
4. **Load Balancing:** Built-in service discovery and load balancing ensure proper traffic distribution.
5. **Self-healing Capability:** Automatically detects and replaces unhealthy containers to ensure high availability.

### **(b) Describe how orchestration tools facilitate automated deployment, scaling, and management of application servers**

Orchestration tools enable automation through:

1. **Declarative Configuration:** Uses YAML files to define desired states, with Kubernetes continuously maintaining them.
2. **Rolling Updates:** Gradually replaces old Pod versions with new ones for zero-downtime deployments.
3. **Blue-Green Deployment:** Runs old and new versions simultaneously for fast switching and rollback.
4. **Auto-scaling:** Automatically adjusts replica count based on resource utilization or metrics.
5. **Configuration Management:** Centralized ConfigMap and Secret management for configuration and sensitive data.

## **2. Explain the difference between a Pod, Deployment, and Service**

Table 1: Comparison of Core Kubernetes Resources

Resource Type	Purpose	Lifecycle	Use Cases
<b>Pod</b>	Smallest deployable unit, contains one or more containers	Ephemeral, IP address may change	Running application containers
<b>Deployment</b>	Manages Pod replicasets, provides update and rollback capabilities	Long-lived, manages Pod lifecycle	Stateless application deployment
<b>Service</b>	Provides stable network endpoint and service discovery	Long-lived	Service exposure and load balancing

### **3. What is a Namespace in Kubernetes? Please list one example.**

A **Namespace** is a mechanism for logical resource isolation within a Kubernetes cluster, used to create separate virtual environments.

#### **Example**

```
apiVersion: v1
kind: Namespace
metadata:
  name: production
```

#### **Use Cases**

- Environment isolation: development, staging, production
- Team isolation: team-a, team-b
- Project isolation: project-x, project-y

### **4. Explain the role of the Kubelet. How do you check the nodes in a Kubernetes cluster? (kubectl command expected).**

#### **4.1 Role of the Kubelet**

- Agent running on each worker node.
- Responsible for Pod lifecycle management.
- Reports node and Pod status to API Server.
- Performs container health checks.

#### **4.2 Commands to Check Nodes**

- View all nodes: `kubectl get nodes`
- Detailed node information: `kubectl describe node <node-name>`
- Check resource usage: `kubectl top nodes`

## 5. What is the difference between ClusterIP, NodePort, and LoadBalancer services?

Table 2: Comparison of Kubernetes Service Types

Service Type	Access Scope	Use Cases	Configuration
<b>ClusterIP</b>	Internal only	Internal service communication	type: ClusterIP
<b>NodePort</b>	NodeIP:Port	Dev/testing, exposing service on every Node's IP	type: NodePort
<b>LoadBalancer</b>	Internet access via Cloud Provider Load Balancer	Cloud production exposure and managed load balancing	type: LoadBalancer

## 6. How do you scale a Deployment to 5 replicas using kubectl?

### Method 1: Direct Scaling (kubectl scale)

- Purpose:** This is the fastest and most direct way to change the replica count of an existing Deployment, suitable for immediate operational adjustments.
- Command:** The command directly targets the Deployment object and patches the `replicas` field.

```
kubectl scale deployment/my-app --replicas=5
```

### Method 2: Edit Deployment (kubectl edit)

- Purpose:** This method is used for interactive editing of the live Deployment resource definition, allowing for modification of multiple fields simultaneously (not just replicas).
- Command:** The command opens the resource's live YAML configuration in a text editor (like Vim or Nano). The user must manually find the `replicas` field under the `spec` section, set its value to 5, and save the file.

```
kubectl edit deployment my-app
```

### Method 3: Apply Updated Configuration File (kubectl apply)

- Purpose:** This is the preferred method for workflows based on **version control (GitOps)**. The desired state is explicitly defined and tracked in a local configuration file.
- Command:** `kubectl` compares the local file state with the cluster state and applies the changes, triggering the scaling event.

```
kubectl apply -f deployment.yaml
```

## 7. How would you update the image of a Deployment without downtime?

Kubernetes achieves zero-downtime updates through the Deployment's default **RollingUpdate** strategy, which gradually replaces the old ReplicaSet with a new one. Based on this mechanism, the image can be updated in several ways.

### Methods Overview

First, the most direct way is using a **command-based update**:

#### Method 1 — Quick Command

```
kubectl blueset image deployment/my-app my-container=nginx:1.21
```

This method triggers a rolling update immediately without modifying any YAML files.

Next, for small on-the-fly adjustments, interactive editing may be used:

#### Method 2 — Live Editing (Not recommended for production)

```
kubectl edit deployment my-app
```

This allows modifying the image field directly in the editor; however, it cannot be version-controlled, making it unsuitable for production environments.

Finally, in real engineering and GitOps workflows, the declarative approach is preferred:

#### Method 3 — Declarative Update (Recommended)

Update the image field in the local `deployment.yaml` and apply the changes:

```
kubectl apply -f deployment.yaml
```

This ensures that configuration changes are traceable and auditable.

### Monitoring and Rollback

After triggering the update, the rollout process can be monitored with:

```
kubectl rollout status deployment/my-app
```

To inspect version history for potential rollback:

```
kubectl rollout bluehistory deployment/my-app
```

If the new version causes issues, the Deployment can be rolled back without downtime:

```
kubectl rollout undo deployment/my-app
```

In summary, all these methods rely on Kubernetes' RollingUpdate mechanism, enabling smooth and uninterrupted

## 8. How do you expose a Deployment to external traffic?

Exposing a Deployment to external clients involves addressing several challenges, such as service discovery, load balancing, NAT, security, and production readiness. The following table summarizes these challenges and corresponding solutions:

Table 3: Challenges and Solutions for Exposing Kubernetes Deployments Externally

Challenge	Description	Solution
<b>Service Discovery</b>	External clients need a stable and reachable endpoint	Stable external IP / DNS via <b>Service or Ingress</b>
<b>Load Balancing</b>	Handling traffic distribution across multiple Pod replicas	Built-in load balancing using Service or Ingress
<b>NAT</b>	Pod IP not externally routable due to Network Address Translation	Use <b>NodePort, LoadBalancer, or Ingress</b>
<b>Security</b>	Protecting the public endpoint from various public access risks	Implement <b>TLS</b> , routing rules, and IP whitelists
<b>Production Readiness</b>	Ensuring the application has High Availability (HA) and scalability	Configure <b>Rolling Updates</b> and zero-downtime deployment

To address these challenges, Kubernetes offers **Ingress** as the recommended solution for HTTP/HTTPS traffic. Ingress provides Layer 7 routing, allowing multiple Services to be exposed via a single external IP or domain. It requires an **Ingress Controller**, such as:

- **Nginx Ingress Controller** — widely used and easy to deploy
- **Traefik** — supports dynamic routing and cloud-native integration
- **HAProxy** — high-performance and highly customizable
- **Istio Ingress** — integrates with service mesh for advanced traffic management
- **AWS ALB Ingress Controller** — managed load balancing for AWS environments

By combining Services and Ingress, Kubernetes can efficiently route external traffic to the appropriate Pods while handling load balancing, security, and high availability.

## 9. How does Kubernetes scheduling decide which node a Pod runs on?

Kubernetes Scheduler determines the optimal node for a Pod to run on, ensuring efficient resource usage, high availability, and adherence to constraints. The scheduler evaluates several factors:

- **Resource requirements:** CPU, memory, and GPU requests/limits are considered to match Pods to suitable nodes.
- **Node selectors (labels):** Pods can specify node labels to control placement on specific nodes.

- **Affinity/Anti-affinity rules:** Influence scheduling based on Pod co-location preferences or avoidance.
- **Taints and tolerations:** Nodes can repel Pods unless they tolerate specific taints.
- **Resource availability and constraints:** Scheduler ensures nodes have enough capacity and obey cluster-level policies.

Overall, the scheduler evaluates these factors to place Pods efficiently while meeting workload and cluster requirements.

## 10. What is the role of Ingress and how does it differ from a Service?

In Kubernetes, an **Ingress** is a resource that manages external access to Services, typically over HTTP or HTTPS. Its primary role is to provide **Layer 7 (application-level) routing**, enabling multiple Services to be exposed under a single IP or domain name. Ingress can perform advanced traffic management functions such as:

- **Domain and path-based routing:** Direct requests to different Services based on hostnames or URL paths.
- **SSL/TLS termination:** Handle encryption at the Ingress level, simplifying certificate management.
- **Virtual hosting:** Serve multiple domains using the same Ingress Controller.
- **Load balancing:** Distribute traffic across multiple backend Pods.
- **Centralized access control:** Apply rules for authentication, whitelisting, or rate-limiting.

To function, an Ingress requires an **Ingress Controller**, which implements the routing rules and handles incoming traffic. Popular controllers include Nginx, Traefik, HAProxy, and cloud-managed options like AWS ALB Ingress Controller.

### Difference Between Ingress and Service

While both Ingress and Service expose Pods to external or internal traffic, they operate at different network layers and serve distinct purposes:

Feature	Service	Ingress
Network Layer	L4 (TCP/UDP)	L7 (HTTP/HTTPS)
Routing	Port-based	Hostname/path-based routing
SSL/TLS	Requires manual configuration	Built-in termination support
External Access	NodePort / LoadBalancer	Single IP with multiple Services via rules
Use Case	Internal or basic external access	Fine-grained external traffic management

Table 4: Comparison between Service and Ingress

In summary, Services provide stable networking for Pods and basic load balancing, while Ingress offers application-layer routing, centralized traffic management, and secure external access for multiple Services.